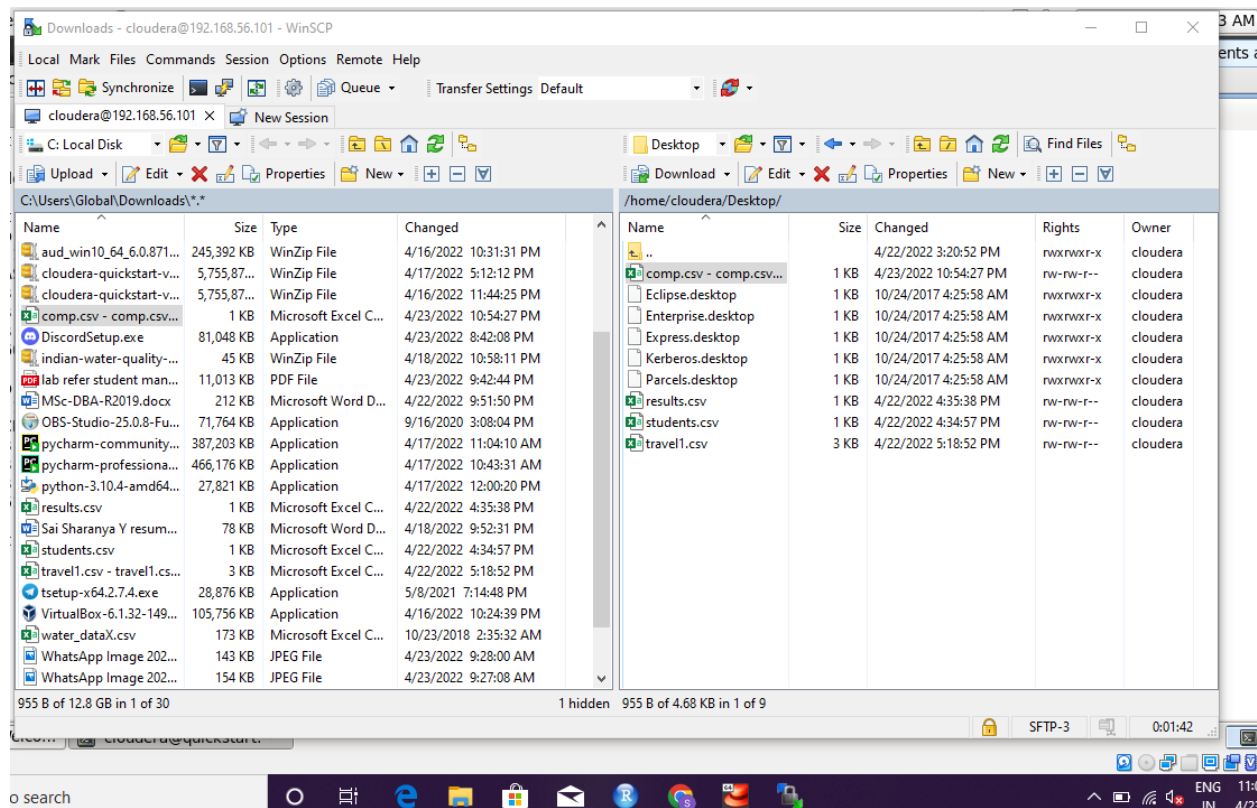


ADVANCE BIG DATA LAB ANALYTICS


EXP - 4

19BTRCR043
SAI SHARANYA Y

To build a script which produces a report listing each Company & State and number of complaints raised by them.



```
[cloudera@quickstart ~]$ hdfs dfs -mkdir /Companydata
[cloudera@quickstart ~]$ hdfs dfs -ls /
Found 9 items
drwxr-xr-x - cloudera supergroup          0 2022-04-23 10:25 /Companydata
drwxr-xr-x - cloudera supergroup          0 2022-04-22 09:18 /TravelLab
drwxrwxrwx - hdfs supergroup              0 2017-10-23 09:15 /benchmarks
drwxr-xr-x - hbase supergroup             0 2022-04-22 02:59 /hbase
drwxr-xr-x - cloudera supergroup          0 2022-04-23 09:46 /salarydata
drwxr-xr-x - solr solr                    0 2017-10-23 09:18 /solr
drwxrwxrwt - hdfs supergroup              0 2022-04-23 09:07 /tmp
drwxr-xr-x - hdfs supergroup              0 2017-10-23 09:17 /user
drwxr-xr-x - hdfs supergroup              0 2017-10-23 09:17 /var
[cloudera@quickstart ~]$ hdfs dfs -put /home/cloudera/Desktop/comp.csv /Companydata
[cloudera@quickstart ~]$ hdfs dfs -ls /Companydata
Found 1 items
-rw-r--r-- 1 cloudera supergroup          955 2022-04-23 10:26 /Companydata/comp.csv
[cloudera@quickstart ~]$
```



```
grunt>
grunt>
grunt>
grunt> compdata = load '/Companydata/comp.csv' using PigStorage(',') as (company_name : chararray, state : chararray, complaint : chararray);
2022-04-23 10:28:52,077 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS
2022-04-23 10:28:52,077 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - mapred.job.tracker is deprecated. Instead, use mapreduce.job.tracker.address
grunt> dump compdata;
2022-04-23 10:29:05,827 [main] INFO org.apache.pig.tools.pigstats.ScriptState - Pig features used in the script: UNKNOWN
2022-04-23 10:29:05,828 [main] INFO org.apache.pig.newplan.logical.optimizer.LogicalPlanOptimizer - {RULES_ENABLED=[AddForEach, ColumnMapKeyPrune, DuplicateForEachCol
umnRewrite, GroupByConstParallelSetter, ImplicitSplitInserter, LimitOptimizer, LoadTypeCastInserter, MergeFilter, MergeForEach, NewPartitionFilterOptimizer, PushDownFo
rEachFlatten, PushUpFilter, SplitFilter, StreamTypeCastInserter], RULES_DISABLED=[FilterLogicExpressionSimplifier, PartitionFilterOptimizer]}
2022-04-23 10:29:05,835 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MRCompiler - File concatenation threshold: 100 optimistic? false
2022-04-23 10:29:05,837 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MultiQueryOptimizer - MR plan size before optimization: 1
2022-04-23 10:29:05,837 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MultiQueryOptimizer - MR plan size after optimization: 1
2022-04-23 10:29:05,869 [main] INFO org.apache.hadoop.yarn.client.RMPProxy - Connecting to ResourceManager at /0.0.0.0:8032
2022-04-23 10:29:05,871 [main] INFO org.apache.pig.tools.pigstats.ScriptState - Pig script settings are added to the job
2022-04-23 10:29:05,873 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler - mapred.job.reduce.markreset.buffer.percent is no
t set, set to default: 0.3
2022-04-23 10:29:06,333 [DataStreamer for file /tmp/temp1844820160/tmp-554985555/hive-hbase-handler-1.1.0-cdh5.13.0.jar] WARN org.apache.hadoop.hdfs.DFSClient - Caught
Exception
java.lang.InterruptedException
    at java.lang.Object.wait(Native Method)
    at java.lang.Thread.join(Thread.java:1281)
    at java.lang.Thread.join(Thread.java:1355)
    at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.closeResponder(DFSOutputStream.java:967)
    at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.endBlock(DFSOutputStream.java:705)
```



```
HadoopVersion 2.6.0-cdh5.13.0 PigVersion 0.12.0-cdh5.13.0 UserId root StartedAt 2022-04-23 10:29:05 FinishedAt 2022-04-23 10:29:36 Features 2022-04-23 10:29:36 UNKNOWN
Success!
Job Stats (time in seconds):
JobId Maps Reduces MaxMapTime MinMapTime AvgMapTime MedianMapTime MaxReduceTime MinReduceTime AvgReduceTime MedianReduceTime Alias F
eature Outputs
job_1650620943421_0012 1 0 8 8 8 8 n/a n/a n/a n/a compdata MAP_ONLY hdfs://quickstart.cloudera:8020
/tmp/temp1844820160/tmp806086029,
Input(s):
Successfully read 19 records (1326 bytes) from: "/Companydata/comp.csv"
Output(s):
Successfully stored 19 records (820 bytes) in: "hdfs://quickstart.cloudera:8020/tmp/temp1844820160/tmp806086029"
Counters:
Total records written : 19
Total bytes written : 820
Spillable Memory Manager spill count : 0
Total bags proactively spilled: 0
Total records proactively spilled: 0
Job DAG:
job_1650620943421_0012
```




```
HadoopVersion  PigVersion  UserId  StartedAt      FinishedAt      Features
2.6.0-cdh5.13.0 0.12.0-cdh5.13.0 root    2022-04-23 10:33:19  2022-04-23 10:34:09  GROUP_BY

Success!


Job Stats (time in seconds):
JobId  Maps    Reduces  MaxMapTime    MinMapTime    AvgMapTime    MedianMapTime    MaxReduceTime    MinReduceTime    AvgReduceTime    MedianReduceTime    Alias F
eature Outputs
job_1650620943421_0013  1      1      15      15      15      15      10      10      10      10      compdata,x  GROUP_BY  hdfs://quickstart.cloudera:8020
/tmp/temp1844820160/tmp-325428902,

Input(s):
Successfully read 19 records (1326 bytes) from: "/Companydata/comp.csv"

Output(s):
Successfully stored 17 records (1445 bytes) in: "hdfs://quickstart.cloudera:8020/tmp/temp1844820160/tmp-325428902"


Counters:
Total records written : 17
Total bytes written : 1445
Spillable Memory Manager spill count : 0
Total bags proactively spilled: 0
Total records proactively spilled: 0

Job DAG:
job_1650620943421_0013
```



```
Job DAG:
job_1650620943421_0013

2022-04-23 10:34:09,286 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - Success!
2022-04-23 10:34:09,287 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS
2022-04-23 10:34:09,287 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - mapreduce.job.tracker is deprecated. Instead, use mapreduce.jobtracker.address
2022-04-23 10:34:09,287 [main] INFO org.apache.pig.data.SchemaTupleBackend - Key [pig.schematuple] was not set... will not generate code.
2022-04-23 10:34:09,294 [main] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input paths to process : 1
2022-04-23 10:34:09,294 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1
(("EQUIFAX, INC."),(("EQUIFAX, INC.",NY)))
(("CITIBANK, N.A."),(("CITIBANK, N.A.",FL)))
(("AES/PHEAA,MA"),(("AES/PHEAA,MA,3475943}))
(("DISCOVER BANK,NV"),(("DISCOVER BANK,NV,1816726}))
(("WINN LAW GROUP, APC"),(("WINN LAW GROUP, APC",CA)))
(("BANK OF AMERICA, NATIONAL ASSOCIATION"),(("BANK OF AMERICA, NATIONAL ASSOCIATION",MD),("BANK OF AMERICA, NATIONAL ASSOCIATION",PA)))
(("Cascade Capital, LLC"),(("Cascade Capital, LLC",FL)))
(("SYNCHRONY FINANCIAL,CA"),(("SYNCHRONY FINANCIAL,CA,3054861}))
(("WELLS FARGO & COMPANY,CA"),(("WELLS FARGO & COMPANY,CA,1527601}))
(("TD BANK US HOLDING COMPANY,NJ"),(("TD BANK US HOLDING COMPANY,NJ,2874325}))
(("BANCO POPULAR DE PUERTO RICO,PR"),(("BANCO POPULAR DE PUERTO RICO,PR,2107177}))
(("CARRINGTON MORTGAGE SERVICES, LLC"),(("CARRINGTON MORTGAGE SERVICES, LLC",ND)))
(("TRANSUNION INTERMEDIATE HOLDINGS, INC."),(("TRANSUNION INTERMEDIATE HOLDINGS, INC.",NC),("TRANSUNION INTERMEDIATE HOLDINGS, INC.",VA)))
(("CAPITAL ONE FINANCIAL CORPORATION,MD"),(("CAPITAL ONE FINANCIAL CORPORATION,MD,3108212}))
(("Experian Information Solutions Inc.,GA"),(("Experian Information Solutions Inc.,GA,2266132}))
(("Experian Information Solutions Inc.,IL"),(("Experian Information Solutions Inc.,IL,3119951}))
(("Experian Information Solutions Inc.,NJ"),(("Experian Information Solutions Inc.,NJ,1509954}))
grunt>
```



```
((Experian Information Solutions Inc.,NJ),{(Experian Information Solutions Inc.,NJ,1509954)})
grunt> y = foreach x generate group, COUNT(compdata.complaint);
grunt> dump y;
2022-04-23 10:42:39,937 [main] INFO org.apache.pig.tools.pigstats.ScriptState - Pig features used in the script: GROUP BY
2022-04-23 10:42:39,938 [main] INFO org.apache.pig.newplan.logical.optimizer.LogicalPlanOptimizer - {RULES_ENABLED=[AddForEach, ColumnMapKeyPrune, DuplicateForEachCol
umnRewrite, GroupByConstParallelSetter, ImplicitSplitInserter, LimitOptimizer, LoadTypeCastInserter, MergeFilter, MergeForEach, NewPartitionFilterOptimizer, PushDownFo
rEachFlatten, PushUpFilter, SplitFilter, StreamTypeCastInserter], RULES_DISABLED=[FilterLogicExpressionSimplifier, PartitionFilterOptimizer]}
2022-04-23 10:42:39,956 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MRCompiler - File concatenation threshold: 100 optimistic? false
2022-04-23 10:42:39,957 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.CombinerOptimizer - Choosing to move algebraic foreach to combiner
2022-04-23 10:42:39,961 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MultiQueryOptimizer - MR plan size before optimization: 1
2022-04-23 10:42:39,961 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MultiQueryOptimizer - MR plan size after optimization: 1
2022-04-23 10:42:40,004 [main] INFO org.apache.hadoop.yarn.client.RMProxy - Connecting to ResourceManager at /0.0.0.0:8032
2022-04-23 10:42:40,006 [main] INFO org.apache.pig.tools.pigstats.ScriptState - Pig script settings are added to the job
2022-04-23 10:42:40,012 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler - mapred.job.reduce.markreset.buffer.percent is no
t set, set to default 0.3
2022-04-23 10:42:40,012 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler - Reduce phase detected, estimating # of required
reducers.
2022-04-23 10:42:40,012 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler - Using reducer estimator: org.apache.pig.backend.
hadoop.executionengine.mapReduceLayer.InputSizeReducerEstimator
2022-04-23 10:42:40,014 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.InputSizeReducerEstimator - BytesPerReducer=1000000000 maxReducers=99
9 totalInputFileSize=955
2022-04-23 10:42:40,014 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler - Setting Parallelism to 1
2022-04-23 10:42:40,356 [DataStreamer for file /tmp/temp1844820160/tmp-2023882259/jdo-api-3.0.1.jar] WARN org.apache.hadoop.hdfs.DFSClient - Caught exception
java.lang.InterruptedException
    at java.lang.Object.wait(Native Method)
    at java.lang.Thread.join(Thread.java:1281)
    at java.lang.Thread.join(Thread.java:1355)
    at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.closeResponder(DFSOutputStream.java:967)
    at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.endBlock(DFSOutputStream.java:705)
```



```
HadoopVersion PigVersion UserId StartedAt FinishedAt Features
2.6.0-cdh5.13.0 0.12.0-cdh5.13.0 root 2022-04-23 10:42:40 2022-04-23 10:43:20 GROUP_BY

Success!


Job Stats (time in seconds):
JobId Maps Reduces MaxMapTime MinMapTime AvgMapTime MedianMapTime MaxReduceTime MinReduceTime AvgReduceTime MedianReduceTime Alias F
eature Outputs
job_1650620943421_0014 1 1 7 7 7 7 7 7 7 7 compdata,x,y GROUP_BY,COMBINER hdfs://quickstart.cloud
era:8020/tmp/temp1844820160/tmp-1698350341,

Input(s):
Successfully read 19 records (1326 bytes) from: "/Companydata/comp.csv"

Output(s):
Successfully stored 17 records (620 bytes) in: "hdfs://quickstart.cloudera:8020/tmp/temp1844820160/tmp-1698350341"

Counters:
Total records written : 17
Total bytes written : 620
Spillable Memory Manager spill count : 0
Total bags proactively spilled: 0
Total records proactively spilled: 0

Job DAG:
job_1650620943421_0014
```



```
2022-04-23 10:43:20,684 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - Success!
2022-04-23 10:43:20,685 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS
2022-04-23 10:43:20,686 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - mapred.job.tracker is deprecated. Instead, use mapreduce.jobtracker.address
2022-04-23 10:43:20,686 [main] INFO org.apache.pig.data.SchemaTupleBackend - Key [pig.schematuple] was not set... will not generate code.
2022-04-23 10:43:20,702 [main] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input paths to process : 1
2022-04-23 10:43:20,702 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1
(("EQUIFAX, INC."),1)
(("CITIBANK, N.A."),1)
(("AES/PHEAR, MA),1)
(("DISCOVER BANK, NV),1)
(("WINN LAW GROUP, APC"),1)
(("BANK OF AMERICA, NATIONAL ASSOCIATION"),2)
(("Cascade Capital, LLC"),1)
(("SYNCHRONY FINANCIAL,CA),1)
(("WELLS FARGO & COMPANY,CA),1)
(("TD BANK US HOLDING COMPANY,NJ),1)
(("BANCO POPULAR DE PUERTO RICO,PR),1)
(("CARRINGTON MORTGAGE SERVICES, LLC"),1)
(("TRANSUNION INTERMEDIATE HOLDINGS, INC."),2)
(("CAPITAL ONE FINANCIAL CORPORATION,MD),1)
((Experian Information Solutions Inc.,GA),1)
((Experian Information Solutions Inc.,IL),1)
((Experian Information Solutions Inc.,NJ),1)
grunt>
```

```

((EXPERIMENTATION) (EXPERIMENTATION) (EXPERIMENTATION) (EXPERIMENTATION) (EXPERIMENTATION)
grunt> z = ORDER y BY $1 DESC;
grunt> dump z;
2022-04-23 10:44:49,587 [main] INFO org.apache.pig.tools.pigstats.ScriptState - Pig features used in the script: GROUP_BY,ORDER_BY
2022-04-23 10:44:49,588 [main] INFO org.apache.pig.newplan.logical.optimizer.LogicalPlanOptimizer - {RULES_ENABLED=[AddForEach, ColumnMapKeyPrune, DuplicateForEachColumnRewrite, GroupByConstParallelSetter, ImplicitSplitInserter, LimitOptimizer, LoadTypeCastInserter, MergeFilter, MergeForEach, NewPartitionFilterOptimizer, PushDownForEachFlatten, PushUpFilter, SplitFilter, StreamTypeCastInserter], RULES_DISABLED=[FilterLogicExpressionSimplifier, PartitionFilterOptimizer]}
2022-04-23 10:44:49,600 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MRCompiler - File concatenation threshold: 100 optimistic? false
2022-04-23 10:44:49,607 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.CombinerOptimizer - Choosing to move algebraic foreach to combiner
2022-04-23 10:44:49,610 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MultiQueryOptimizer - MR plan size before optimization: 3
2022-04-23 10:44:49,610 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MultiQueryOptimizer - MR plan size after optimization: 3
2022-04-23 10:44:49,637 [main] INFO org.apache.hadoop.yarn.client.RMProxy - Connecting to ResourceManager at /0.0.0.0:8032
2022-04-23 10:44:49,639 [main] INFO org.apache.pig.tools.pigstats.ScriptState - Pig script settings are added to the job
2022-04-23 10:44:49,661 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler - mapred.job.reduce.markreset.buffer.percent is not set, set to default 0.3
2022-04-23 10:44:49,664 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler - Reduce phase detected, estimating # of required reducers.
2022-04-23 10:44:49,664 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler - Using reducer estimator: org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.InputSizeReducerEstimator
2022-04-23 10:44:49,666 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.InputSizeReducerEstimator - BytesPerReducer=1000000000 maxReducers=999 totalInputFileSize=955
2022-04-23 10:44:49,666 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler - Setting Parallelism to 1
2022-04-23 10:44:50,013 [DataStreamer for file /tmp/temp1844820160/tmp589044177/hive-hbase-handler-1.1.0-cdh5.13.0.jar] WARN org.apache.hadoop.hdfs.DFSClient - Caught exception
java.lang.InterruptedException
    at java.lang.Object.wait(Native Method)
    at java.lang.Thread.join(Thread.java:1281)

```



```

Success!

Job Stats (time in seconds):
JobId  Maps  Reduces MaxMapTime  MinMapTime  AvgMapTime  MedianMapTime  MaxReduceTime  MinReduceTime  AvgReduceTime  MedianReduceTime  Alias F
eature Outputs
job_1650620943421_0015  1  1  7  7  7  7  8  8  8  8  compdata,x,y  GROUP_BY,COMBINER
job_1650620943421_0016  1  1  6  6  6  6  8  8  8  8  z  SAMPLER
job_1650620943421_0017  1  1  7  7  7  7  8  8  8  8  z  ORDER_BY
p1844820160/tmp-1261696629,
                                hdfs://quickstart.cloudera:8020/tmp/tem

Input(s):
Successfully read 19 records (1326 bytes) from: "/Companydata/comp.csv"

Output(s):
Successfully stored 17 records (620 bytes) in: "hdfs://quickstart.cloudera:8020/tmp/temp1844820160/tmp-1261696629"

Counters:
Total records written : 17
Total bytes written : 620
Spillable Memory Manager spill count : 0
Total bags proactively spilled: 0
Total records proactively spilled: 0

Job DAG:
job_1650620943421_0015  ->  job_1650620943421_0016,
job_1650620943421_0016  ->  job_1650620943421_0017,
job_1650620943421_0017

```




```
root@quickstart:~
at org.apache.hadoop.util.RunJar.run(RunJar.java:221)
at org.apache.hadoop.util.RunJar.main(RunJar.java:136)
2022-04-23 10:46:57,298 [main] INFO org.apache.hadoop.mapred.ClientServiceDelegate - Application state is completed. FinalApplicationStatus=SUCCEEDED. Redirecting to
job history server
2022-04-23 10:46:57,583 [main] INFO org.apache.hadoop.mapred.ClientServiceDelegate - Application state is completed. FinalApplicationStatus=SUCCEEDED. Redirecting to
job history server
2022-04-23 10:46:57,760 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - Success!
2022-04-23 10:46:57,761 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS
2022-04-23 10:46:57,761 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - mapred.job.tracker is deprecated. Instead, use mapreduce.jobtracker.address
2022-04-23 10:46:57,761 [main] INFO org.apache.pig.data.SchemaTupleBackend - Key [pig.schematuple] was not set... will not generate code.
2022-04-23 10:46:57,766 [main] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input paths to process : 1
2022-04-23 10:46:57,766 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1
(("TRANSSION INTERMEDIATE HOLDINGS, INC."),2)
(("BANK OF AMERICA, NATIONAL ASSOCIATION"),2)
(("Experian Information Solutions Inc.,NJ"),1)
(("Experian Information Solutions Inc.,IL"),1)
(("Experian Information Solutions Inc.,GA"),1)
(("CAPITAL ONE FINANCIAL CORPORATION,MD"),1)
(("CARRINGTON MORTGAGE SERVICES, LLC"),1)
(("BANCO POPULAR DE PUERTO RICO,PR"),1)
(("TD BANK US HOLDING COMPANY,NJ"),1)
(("WELLS FARGO & COMPANY,CA"),1)
(("SYNCHRONY FINANCIAL,CA"),1)
(("Cascade Capital, LLC"),1)
(("WINN LAW GROUP, APC"),1)
(("DISCOVER BANK,NV"),1)
(("AES/PHEAA,MA"),1)
(("CITIBANK, N.A."),1)
(("EQUIFAX, INC."),1)
grunt>

[Cloudera Live : Welco... root@quickstart:~
```

```
grunt> store z into '/Companydata/required';
2022-04-23 10:49:25,760 [main] INFO org.apache.pig.tools.pigstats.ScriptState - Pig features used in the script: GROUP BY,ORDER BY
2022-04-23 10:49:25,761 [main] INFO org.apache.pig.newplan.logical.optimizer.LogicalPlanOptimizer - {RULES_ENABLED=[AddForEach, ColumnMapKeyPrune, DuplicateForEachCol
umnRewrite, GroupByConstParallelSetter, ImplicitSplitInserter, LimitOptimizer, LoadTypeCastInserter, MergeFilter, MergeForEach, NewPartitionFilterOptimizer, PushDownFo
rEachFlatten, PushUpFilter, SplitFilter, StreamTypeCastInserter], RULES_DISABLED=[FilterLogicExpressionSimplifier, PartitionFilterOptimizer]}
2022-04-23 10:49:25,769 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MRCompiler - File concatenation threshold: 100 optimistic? false
2022-04-23 10:49:25,776 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.CombinerOptimizer - Choosing to move algebraic foreach to combiner
2022-04-23 10:49:25,779 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MultiQueryOptimizer - MR plan size before optimization: 3
2022-04-23 10:49:25,780 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MultiQueryOptimizer - MR plan size after optimization: 3
2022-04-23 10:49:25,889 [main] INFO org.apache.hadoop.yarn.client.RMProxy - Connecting to ResourceManager at /0.0.0.0:8032
2022-04-23 10:49:25,811 [main] INFO org.apache.pig.tools.pigstats.ScriptState - Pig script settings are added to the job
2022-04-23 10:49:25,818 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler - mapred.job.reduce.markreset.buffer.percent is no
t set, set to default 0.3
2022-04-23 10:49:25,818 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler - Reduce phase detected, estimating # of required
reducers.
2022-04-23 10:49:25,818 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler - Using reducer estimator: org.apache.pig.backend.
hadoop.executionengine.mapReduceLayer.InputSizeReducerEstimator
2022-04-23 10:49:25,820 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.InputSizeReducerEstimator - BytesPerReducer=1000000000 maxReducers=99
9 totalInputFileSize=955
2022-04-23 10:49:25,820 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler - Setting Parallelism to 1
2022-04-23 10:49:26,123 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler - creating jar file Job6946078999961941446.jar
2022-04-23 10:49:30,096 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler - jar file Job6946078999961941446.jar created
2022-04-23 10:49:30,102 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler - Setting up single store job
2022-04-23 10:49:30,103 [main] INFO org.apache.pig.data.SchemaTupleFrontend - Key [pig.schematuple] is false, will not generate code.
2022-04-23 10:49:30,103 [main] INFO org.apache.pig.data.SchemaTupleFrontend - Starting process to move generated code to distributed cache
2022-04-23 10:49:30,103 [main] INFO org.apache.pig.data.SchemaTupleFrontend - Setting key [pig.schematuple.classes] with classes to deserialize []
2022-04-23 10:49:30,119 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - 1 map-reduce job(s) waiting for submission.
2022-04-23 10:49:30,119 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - mapred.job.tracker is deprecated. Instead, use mapreduce.jobtracker.address

[Cloudera Live : Welco... root@quickstart:~
```